

CRS AI-Provenance Classifiers — Final Results

Generated 2026-05-20T06:37:01.711121+00:00

Backbone: microsoft/deberta-v3-large LoRA rank 16 3 epochs AI chunks pooled & randomly split; HUMAN temporally split; no OOD.

Model A — 6-class (with HUMAN)



Train dist: GPT:8599 HUMAN:8000 DEEPSEEK:8539
GEMINI:8599 CLAUDE:8500 GROK:8500

class	P	R	F1	n
HUMAN	0.273	0.374	0.316	500
CLAUDE	0.751	0.590	0.661	500
GPT	0.379	0.648	0.478	500
GROK	0.367	0.184	0.245	500
GEMINI	0.944	0.870	0.905	500
DEEPSEEK	0.407	0.290	0.339	500
macro	—	0.493	0.491	—
accuracy				0.493
ECE				0.0571
HUMAN→AI FPR				0.626

Model B — 5-class (AI only)

Train dist: GPT:8599 DEEPSEEK:8539 GEMINI:8599
CLAUDE:8500 GROK:8500

class	P	R	F1	n
CLAUDE	0.589	0.442	0.505	500
GPT	0.408	0.730	0.523	500
GROK	0.315	0.364	0.338	500
GEMINI	0.912	0.812	0.859	500
DEEPSEEK	0.361	0.150	0.212	500
macro	—	0.500	0.488	—
accuracy				0.500
ECE				0.0353



Model A checkpoint: /home/srinivas-harish/nyu-ai-detection-pipeline/workspace/models/checkpoints/crs_final_6class/seed42/
Model B checkpoint: /home/srinivas-harish/nyu-ai-detection-pipeline/workspace/models/checkpoints/crs_final_5class/seed42/
Fit wall: A=9272s, B=8679s.